

Internet Technology Fundamentals

by Ted Kosan

Part of The Professor And Pat series
(professorandpat.org)

Copyright © 20014 by Ted Kosan

This work is licensed under the Creative Commons
Attribution-NonCommercial-NoDerivs 3.0 License. To view a copy of
this license, visit <http://creativecommons.org/licenses/by-nc-nd/3.0/>

Table of Contents

Internet Technology Fundamentals.....	1
1 How does one computer communicate with another computer?.....	3
2 How do multiple computers communicate with each other?.....	4
3 The TCP/IP protocol suite.....	6
4 Clients and servers.....	8
5 DHCP.....	8
6 DNS.....	9
11 Processes and ports.....	12
12 Well known ports, registered ports, and dynamic ports.....	15
12.2 Well known ports (0 - 1023).....	16
12.5 Registered ports (1024 - 49151).....	18
12.6 Dynamic/private ports (49152 - 65535).....	18
13 The SSH (Secure SHell) service	18
14 Using SSH to remotely log into a machine.....	20
15 Using scp to copy files between machines on the network.....	21

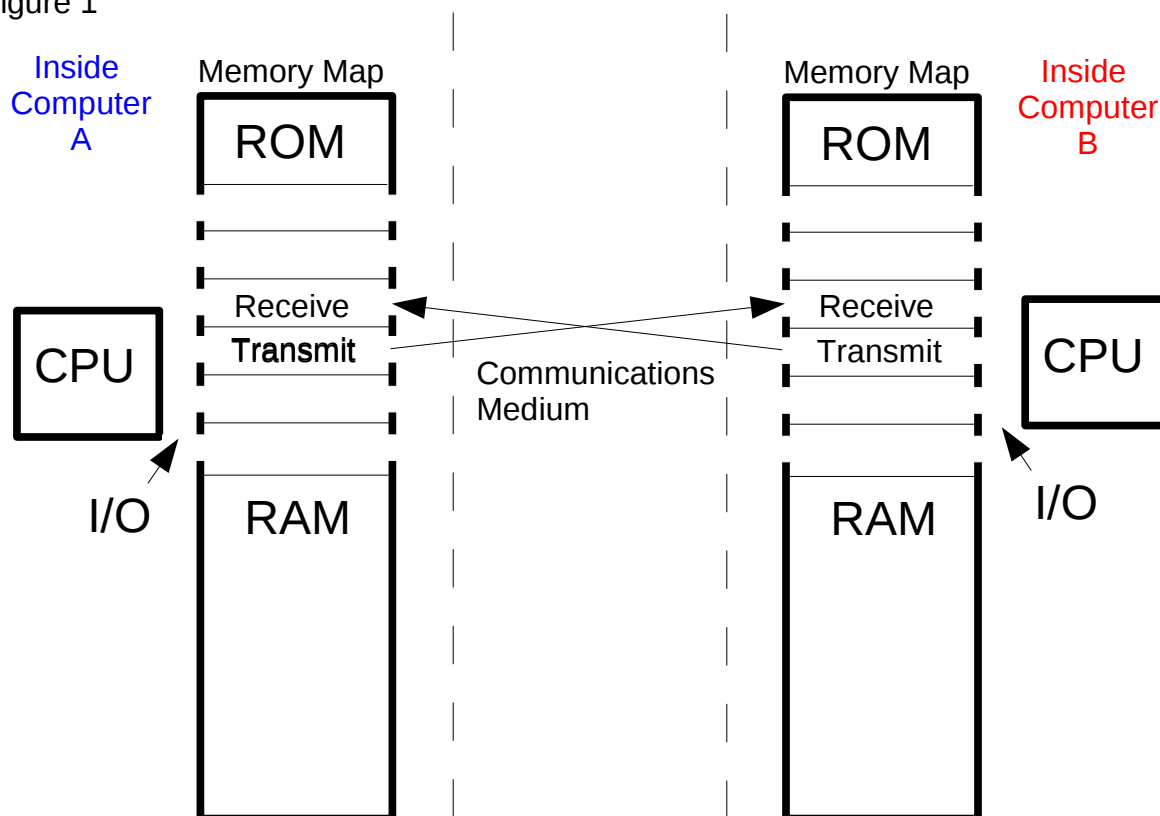
1 How does one computer communicate with another computer?

1.1 The Internet is currently one of the most important technologies of our civilization, and its importance will only increase in the future. The Internet is expanding so quickly that projections show almost all computing devices will eventually be connected to it. Therefore, it is essential for anyone who has a desire to become deeply involved with computers to understand how Internet-related technologies work.

1.2 Understanding the history of how the Internet was created is also important, but we will not be discussing this history here because it has been documented elsewhere. I highly recommend that you do an Internet search on the history of the Internet and read some of the articles you find. I assure you that it will be an excellent investment of your time.

1.3 We are going to approach the topic of how Internet-related technologies work by going back to the model of how a computer works (which was discussed in the “Computer Systems: Gateways To Cyberspace” book) and extending it to show how two computers can be connected to each other so

Figure 1



that they can communicate.

1.4 In the “Computer Systems: Gateways To Cyberspace” book, we discussed how a computer's memory map holds three kinds of memory (RAM, ROM and I/O) and that I/O memory was how a computer communicated with devices outside itself, like keyboards and printers. Since most computers are external to each other, I/O locations are also the mechanism that is used to allow one computer to communicate with another computer. Figure 1 shows two computers which are labeled "A" and "B". Both computers are using one I/O location as a **transmit location** and one I/O location as a **receive location**.

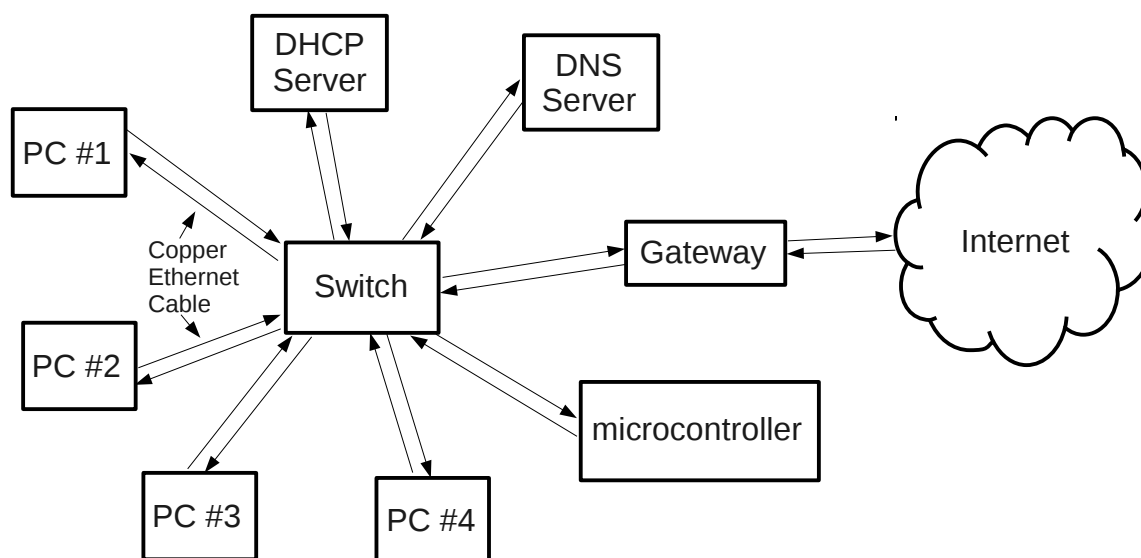
1.5 Computer "A's" **transmit** location is attached to computer "B's" **receive** location (and computer "B's" **transmit** location is attached to computer "A's" **receive** location) using a **communications medium**. A **communications medium** is something that is able to carry information from one point to another. In the case of the model in Figure 1, when computer "A's" CPU places a number in computer "A's" transmit location, the communications medium (represented by the right-pointing arrow) copies this number to computer "B's" receive location. The communications medium represented by the left-pointing arrow will copy numbers that are placed into "B's" transmit location into "A's" receive location.

1.6 There are many kinds of communications mediums, including copper wires, fiber optic cables, and wireless radio signals. One of the most popular communications mediums for networking PCs is called **Ethernet** and most PCs which are sold today include an Ethernet interface. No matter which communications medium a device uses, however, they all perform the same task of copying numbers from the I/O memory locations of one computer to the I/O locations of another computer.

2 How do multiple computers communicate with each other?

2.1 When only two computers need to communicate, the situation is simple because the information that leaves one computer is sent to the other computer and vice versa. But what about the situation where multiple computers need to communicate with each other? There are a number of ways to solve this problem, and one of the more common ways is shown in Figure 2:

Figure 2 Local Area Network (LAN)



2.2 Figure 2 shows multiple computers connected to what is called a **Local Area Network** or **LAN**. A **LAN** consists of multiple computers that are physically close to each other (usually in the same room or in the same building) and attached to each other using some kind of communications medium. In Figure 2, the computers are attached to a device called a **switch** with copper Ethernet cables.

2.3 Computers on a network communicate with each other using **messages**, and sending a message is similar to sending a letter through the mail. The purpose of a **switch** is to look at each message that is sent into it, determine which computer the message is being sent to, and then sending the message to that computer.

2.4 There is a problem with the model in Figure 2, however, because the names that are associated with each computer on the network would not be suitable for uniquely identifying them if their numbers would be increased into the hundreds or thousands. Beyond this, the cloud on the right side of the figure represents the Internet, and the millions of computers (which are also called **hosts**) that are currently attached to it. Messages can also be sent to these computers and received from them, but only if each computer on the Internet is uniquely identified in some way. Beyond this, rules for how the messages are to be exchanged must also exist.

71 3 The TCP/IP protocol suite


72 3.1 Two problems that needed to be solved before the Internet could be
73 created were that each computer needed to be uniquely identified, and
74 communications rules (also called **protocols**) needed to be developed which
75 determined how the messages were to be exchanged. With respect to the
76 Internet, a **protocol** can be defined as "**a set of rules that define an exact
77 format for communication between systems.**"
78 (www.unitedyellowpages.com/internet/terminology.html). When a number
79 of protocols are used together, they are called a **protocol suite**.

80 3.2 The protocol suite that was developed for the Internet is called **TCP/IP**,
81 and its name is a combination of the names of the two most heavily used
82 protocols in the suite (**TCP** stands for **Transmission Control Protocol** and
83 **IP** stands for **Internet Protocol**). The **Internet Protocol** defines a way to
84 uniquely identify computers on the Internet using an addressing system. IP
85 version 4 (**IPv4**), which is currently the most widely used version of the IP
86 protocol, consists of **4 numbers between 0 and 255 separated from
87 each other by a dot**. Examples of IP address include 207.21.94.50,
88 54.3.59.2, and 204.74.99.100. All the addresses from 0.0.0.0 to
89 255.255.255.255 create an **address space** that contains 4,294,967,296
90 addresses.

91 3.3 IP version 6 (**IPv6**) is the newest version of the IP protocol and it has an
92 address space that contains
93 340,282,366,920,938,463,463,374,607,431,768,211,456 addresses! The
94 transition from IPv4 to IPv6 has begun, but it is moving slowly. A large
95 number of hosts on the Internet will continue to use the IPv4 protocol for a
96 long time and therefore IPv4 is what we will use in this document.

97 3.4 Figure 3 contains the same model of a network that was shown in Figure 2
98 but with **IPv4** addresses assigned to each computer:

Figure 3 Local Area Network (LAN) with IPv4 addresses



3.5 If PC #3 needed to send a message to PC #4, the IP address of PC #4 (which is 207.21.94.214) would be placed into the message. The IP address of the sender (207.21.94.72) is also placed into the message in case PC #4 wanted to send a reply (this is similar to placing a return address on a letter). PC #3 will then send the message to the switch, the switch will look at the message's destination address and then pass the message to PC #4.

3.6 If one of the computers on this local network needs to send a message to a computer which is not on the LAN, then the message is sent to the **gateway** computer, and the gateway will then route the message to the Internet.

3.7 During the base Gentoo installation procedure (and also the GUI installation procedure) you have been logging directly into the superuser's account because you needed the extra privileges that this account possessed. **Normally**, however, you should log into a Linux system using your **user account** and then log into the **superuser** account (using the **su** command) only when you need the superuser account's extra privileges.

3.8 Log into the computer you installed Gentoo Linux on **using your user account** and let's determine its **IP address**. Execute an **ifconfig** command in order to determine the IP address of your machine:

```

kosan1 / # ifconfig
eth0      Link encap:Ethernet  HWaddr 00:16:D4:0B:1A:3A
          inet addr:206.21.94.132  Bcast:206.21.94.255  Mask:255.255.255.0

```

```
120      UP BROADCAST NOTRAILERS RUNNING MULTICAST  MTU:1500  Metric:1
121      RX packets:26727 errors:0 dropped:0 overruns:0 frame:0
122      TX packets:22929 errors:0 dropped:0 overruns:0 carrier:0
123      collisions:0 txqueuelen:1000
124      RX bytes:26221365 (25.0 Mb)  TX bytes:4167216 (3.9 Mb)
125      Interrupt:18
```

```
126  lo      Link encap:Local Loopback
127          inet addr:127.0.0.1  Mask:255.0.0.0
128          UP LOOPBACK RUNNING  MTU:16436  Metric:1
129          RX packets:14 errors:0 dropped:0 overruns:0 frame:0
130          TX packets:14 errors:0 dropped:0 overruns:0 carrier:0
131          collisions:0 txqueuelen:0
132          RX bytes:756 (756.0 b)  TX bytes:756 (756.0 b)
```

133 3.9 The IP address of my machine is 206.21.94.132 and it was obtained from
134 the DHCP (Dynamic Host Configuration Server) which is on the LAN that my
135 machine is attached to.

136 4 Clients and servers

137 4.1 On LANs and on the Internet, there are a number of ways for
138 communications between computers to be organized and these
139 organizations are often called **architectures**. One architecture is called
140 **Peer-to-Peer** (P2P), and it treats computers on the network as equals that
141 exchange information with each other. An example of a P2P application is
142 instant messaging.

143 4.2 Another architecture that is used with networked computers is called
144 **Client-Server**. With a Client-Server architecture, a **server** is a computer
145 that accepts requests from other computers on the network, performs the
146 work that was requested, and returns the results of the work to the
147 requester. A **client** is a computer that sends a request to a server, receives
148 a response, and then makes use of the information that was contained in the
149 response.

150 4.3 In the LAN shown in Figure 3, there are two servers (a DHCP server and a
151 DNS server) and five clients. The servers will be discussed in the next two
152 sections.

153 5 DHCP

154 5.1 **DHCP** stands for **Dynamic Host Configuration Protocol** and its
155 purpose is to allow computers on a LAN to automatically be configured
156 when they are booted up with the information they need to access the

158 network. This information includes an **IP address**, the **address of the**
159 **gateway**, and the **address of a DNS server**. We have already discussed
160 what an IP address is and what a gateway is. DNS servers will be covered
161 in the next section.

162 5.2 What you might be wondering at this point is how a computer that doesn't
163 have an IP address yet (because it is booting up) is able to use the network
164 to contact the DHCP server to obtain an IP address. This problem is solved
165 by having the booting computer send a DHCP **broadcast** message to the
166 LAN. Broadcast messages are not sent to any specific machine on a LAN.
167 Instead, broadcast messages are sent to the LAN as a whole and all the
168 computers that are on the LAN receive the message.

169 5.3 If a DHCP request message is broadcast to the LAN, the DHCP server will
170 receive the request at the same time that the rest of the computers do. The
171 other computers will read the contents of the message, see that it contains a
172 DHCP request, and then they will ignore it. The DHCP server, however, will
173 read the contents of the message, see that the message was meant for it,
174 and send DHCP configuration information back to the sender.

175 5.4 Earlier we saw how the **ifconfig** command could be used to list the IP
176 address that the DHCP server gave your machine. You can also determine
177 what address the DHCP server gave for the **gateway** on your network by
178 using the **netstat** command:

```
179 kasan1 / # netstat -nr
180 Kernel IP routing table
181 Destination      Gateway          Genmask         Flags   MSS Window  irtt Iface
182 0.0.0.0           206.21.94.1     0.0.0.0         UG        0 0          0 enp0s3
183 127.0.0.0         0.0.0.0         255.0.0.0       U          0 0          0 lo
184 206.21.94.0       0.0.0.0         255.255.255.0   U          0 0          0 enp0s3
```

185 5.5 This listing shows that the address for the **gateway** on the network I am
186 using is **206.21.94.1**.

187 6 DNS

188 6.1 Each of the millions of computers on the Internet can be accessed using
189 their IP addresses. For example, the IP address of the server that contains
190 the main Shawnee State University website is **146.85.50.73**. You can
191 access this website by launching a web browser and then entering
192 **http://146.85.50.73/** in the URL bar.

193 6.2 It is difficult for humans to remember numerous numbers, however, so a
194 **system for associating names with IP address numbers** was created for

the Internet. The name of the system is **DNS**, and it stands for **Domain Name System**. A name that is associated with one or more IP address is called a **domain name**, and a **domain name** that has a given machine's **hostname** at its beginning (and a period at its end) is called a **fully qualified domain name**. Examples of domain names are:

gentoo.org
yahoo.com
sourceforge.net
google.com
java.net
wikipedia.com

6.3 Examples of fully qualified domain names are:

kiwi.gentoo.org.
loon.gentoo.org.
wren.gentoo.org.

6.4 DNS is implemented as a large database that is distributed across the whole Internet. Domain names need to be registered with a **domain name registry** organization before they will be entered into the DNS system. Examples of domain name registry companies include godaddy.com, networksolutions.com, and register.com.

6.5 The DNS server on the LAN in Figure 3 has three functions. The first function is to accept messages that contain **domain names** from clients and to return the **IP address** that are associated with these names. When a user types in a domain name like **shawnee.edu** into a browser's URL bar, the browser cannot contact a the server yet because it does not know its IP address. The operating system that the browser is running on will therefore send the domain name to the DNS server (using the DNS server's IP address that it obtained through DHCP), and the DNS server will respond with one or more IP address that are associated with the **shawnee.edu** domain name. The system will then use one of these IP address to contact the Shawnee State University server.

6.6 A Gentoo Linux system holds DNS server IP addresses in the **/etc/resolv.conf** file. Take a moment to look inside this file on your system to see which DNS servers your system has been configured with.

6.7 The second function that a local DNS server has is to **define** the **domain name** to **IP address** mappings for the machines on the local network. If a remote computer on the Internet wants to know the IP address for a

232 machine on the local network, and its DNS server does not know the
233 mapping, the remote DNS server will contact the local **authoritative** DNS
234 server to ask what the mapping is. The remote DNS server will then
235 remember this mapping for a certain time in case machines on the remote
236 network need to know the mapping in the future.

237 6.8 The third function that a DNS server has is to take messages that contain
238 **IP addresses** and return the **domain names** that are associated with these
239 addresses.

240 6.9 Now that you know what a DNS server does, lets **emerge** a program that
241 will allow us to **query our local DNS servers**. The program is called **dig**,
242 and it is contained in the **bind-tools package**. The superuser account must
243 be used to execute the emerge command. The superuser account is entered
244 from a user account using the **su** (superuser) command:

```
7  
245 8 tkosan@kosan1 / $ su  
246 9 Password: *****  
247 10 kosan1 / #
```

248 10.1 Notice that the command prompt indicates that a person is in their user
249 account by showing the user's name and the machine's name (in green)
250 along with a dollar '\$'. After entering the superuser's account using the **su**
251 command, only the machine's name is shown (in red) along with a number
252 sign '#' instead of a dollar '\$' sign.

253 10.2 Now that you are in the superuser's account, emerge the bind-tools
254 program:

```
255 kosan1 / # emerge bind-tools
```

256 10.3 After the package has been emerged, lets ask **dig** what IP addresses are
257 associated with the **shawnee.edu** domain name:

```
258 kosan1 / # dig shawnee.edu  
259 ; <<>> DiG 9.9.4 <<>> shawnee.edu  
260 ;; global options: +cmd  
261 ;; Got answer:  
262 ;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 22307  
263 ;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 0  
264  
265 ;; QUESTION SECTION:  
266 ;shawnee.edu.                IN      A  
267  
268 ;; ANSWER SECTION:  
269 shawnee.edu.                 72159 IN      A      146.85.50.73  
270
```

```
271 ;; Query time: 15 msec
272 ;; SERVER: 206.21.94.6#53(206.21.94.6)
273 ;; WHEN: Sun Apr 06 20:29:53 EST 2014
274 ;; MSG SIZE rcvd: 45
```

275 10.4 The dig program indicates that one IP address is associated the
276 **shawnee.edu** domain name. If we want to see what fully qualified domain
277 names are associated with this IP address, we can have **dig** find out for us
278 by executing a **dig -x 146.85.50.73** command:

```
279 kosan1 / # dig -x 146.85.50.73

280 ; <<>> DiG 9.9.4 <<>> -x 146.85.50.73
281 ;; global options: +cmd
282 ;; Got answer:
283 ;; ->HEADER<<- opcode: QUERY, status: NOERROR, id: 53876
284 ;; flags: qr rd ra; QUERY: 1, ANSWER: 3, AUTHORITY: 0, ADDITIONAL: 0
285
286 ;; QUESTION SECTION:
287 73.50.85.146.in-addr.arpa. IN PTR
288
289 ;; ANSWER SECTION:
290 73.50.85.146.in-addr.arpa. 86400 IN PTR shawnee.edu.
291 73.50.85.146.in-addr.arpa. 86400 IN PTR www.shawnee.edu.
292 73.50.85.146.in-addr.arpa. 86400 IN PTR omniupdate.shawnee.edu.
293
294 ;; Query time: 17 msec
295 ;; SERVER: 206.21.94.6#53(206.21.94.6)
296 ;; WHEN: Sun Apr 06 20:30:42 EST 2014
297 ;; MSG SIZE rcvd: 111
```

298 11 Processes and ports

299 11.1 Now that we have discussed some of the more important technologies
300 that are related to the Internet, it is time talk about what happens when IP
301 messages (referred to as messages from now on) arrive at a computer, and
302 what generates messages before they are sent from a computer.

303 11.2 Almost all modern personal computers can have multiple programs
304 running on them concurrently. Here is a list of programs that may be
305 running concurrently on a typical user's computer:

- 306 - Web browser.
- 307 - Instant message client.
- 308 - Word processor.
- 309 - File download utility.
- 310 - Audio file player.
- 311 - Computer game.

312 11.3 In Linux, running programs are called **processes**, and a list of all the
 313 **processes** that are currently running on a Linux system can be obtained by
 314 executing a **ps -e** command:

```

315 kosan1 / # ps -e
316   PID TTY          TIME CMD
317     1 ?            00:00:00 init
318     2 ?            00:00:00 ksoftirqd/0
319     3 ?            00:00:00 events/0
320     4 ?            00:00:00 khelper
321     5 ?            00:00:00 kthread
322     8 ?            00:00:00 kblockd/0
323     9 ?            00:00:00 kacpid
324    55 ?            00:00:00 kseriod
325    58 ?            00:00:00 khubd
326   145 ?            00:00:00 pdflush
327   146 ?            00:00:00 pdflush
328   147 ?            00:00:00 kswapd0
329   148 ?            00:00:00 aio/0
330   149 ?            00:00:00 cifsoplockd
331   150 ?            00:00:00 cifsnotifyd
332   753 ?            00:00:00 kpsmouse
333   814 ?            00:00:00 kjournald
334   925 ?            00:00:00 udevd
335  3532 ?            00:00:00 syslog-ng
336  3947 ?            00:00:00 dhclient
337  4168 ?            00:00:00 cron
338  4245 tty1          00:00:00 login
339  4246 tty2          00:00:00 agetty
340  4247 tty3          00:00:00 agetty
341  4258 tty4          00:00:00 agetty
342  4259 tty5          00:00:00 agetty
343  4260 tty6          00:00:00 agetty
344  4285 tty1          00:00:00 bash
345  4289 ?            00:00:00 sshd
346  4292 pts/0        00:00:00 bash
347  4296 tty1          00:00:00 startx
348  4312 tty1          00:00:00 xinit
349  4313 tty7          00:00:12 X
350  4317 tty1          00:00:00 jwm
351  4349 tty1          00:00:00 sh
352  4350 tty1          00:00:00 xload
353  4352 tty1          00:00:00 rxvt
354  4353 pts/1        00:00:00 bash
355  4356 tty1          00:00:00 mozilla-launche
356  4365 tty1          00:00:12 firefox-bin
357  4381 tty1          00:00:00 soffice
358  4390 tty1          00:00:12 soffice.bin
359  4496 pts/0        00:00:00 ps

```

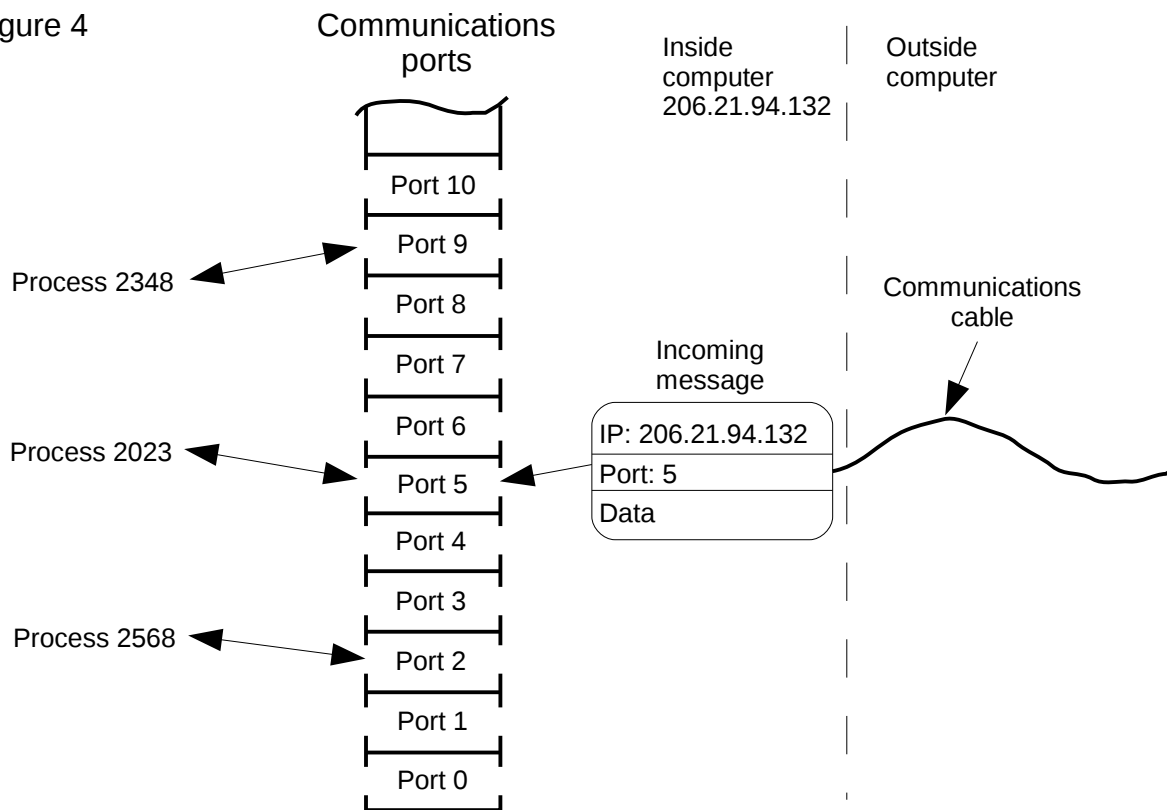
360 11.4 If you look towards the bottom of this list you can see that my computer
 361 is currently running X Windows, the jwm window manager, the rxvt terminal
 362 emulator, a bash shell (that is attached to the terminal emulator), the firefox

363 browser, and the open office word processor. Notice that the **ps** command
 364 included itself in the list because it was running at the moment that the list
 365 was created.

366 11.5 There are four columns in this listing. Each process is given a unique
 367 **Process ID** (PID) number when the process is created, and these numbers
 368 are listed in the **PID** column. The **TTY** column indicates whether or not a
 369 process is attached to a terminal, and if it is, what terminal it is attached to.
 370 The **TIME** column indicates how much CPU time the process has used so far
 371 in hours, minutes and seconds .

372 11.6 When a message arrives at a computer from the network, the computer
 373 must decide which process to give the message to. The way that the TCP/IP
 374 protocol solves this problem is with software-based communications **ports**.

Figure 4



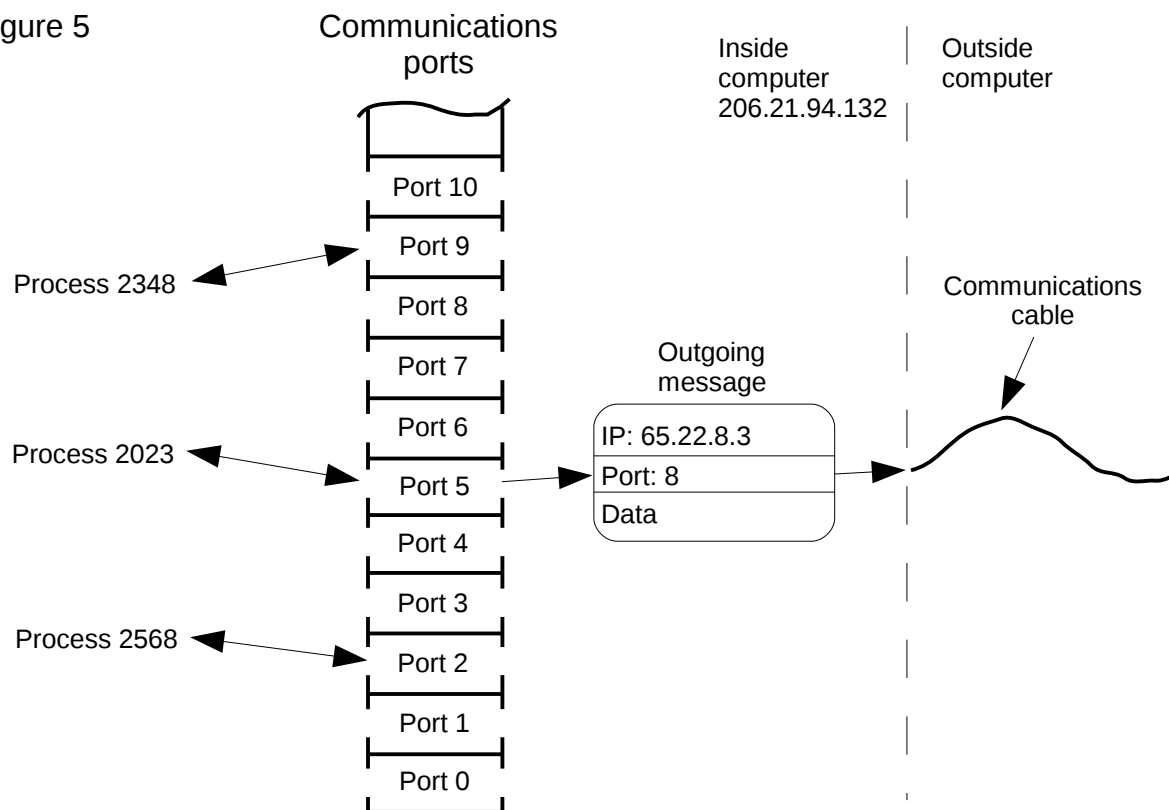
375 11.7 Figure 4 shows the inside and the outside of a computer that is
 376 connected to a network and which has an IP address of **206.21.94.132**.
 377 The communications ports are placed between the processes that are
 378 running on the left and the network connection on the right. Each port is
 379 given a unique number with the lowest port number being **0** and the highest

377 port number being **65535**. Each message that arrives from the network has
378 a port number included in it so that the system knows which port to send
379 the message to.

380 11.8 In Figure 4, a message that has **port 5** as its destination port has arrived
381 from the network and therefore the system will place this message into **port**
382 **5**. **Process 2023** has been bound to **port 5**. When the system sends the
383 message to this port, **process 2023** will take the message and then do
384 something with the information it contains.

385 11.9 Figure 5 shows a message from process **2023** being sent to another
386 computer on the network which has an IP address of **65.22.8.3**. When this
387 messages arrives at the destination computer, it will place the message into
388 it's **port 8**. If there is a process on the destination computer that is bound to
389 **port 8**, it will receive this message.

Figure 5



390 12 Well known ports, registered ports, and dynamic ports

391 12.1 Now that you know what ports are and how processes are bound to
392 them, you may be wondering how people determine which processes should

be bound to which ports. An organization called **IANA** (Internet Assigned Numbers Authority) is responsible for various number schemes associated with the Internet, and one of them is the TCP/IP port scheme. IANA has divided the **0 - 65535** port range into the following three address blocks:

0 - 1023 -> Well Known Ports.

1024 - 49151 -> Registered Ports.

49152 - 65535 -> Dynamic and or Private Ports.

12.2 Well known ports (0 - 1023)

12.2.1 A list is maintained by IANA which indicates which kinds of programs are usually bound to specific port numbers in this range. For example, **web servers** are bound to **port 80**, **SSH (secure shell) servers** are bound to **port 22**, **FTP (File Transfer Protocol) servers** are bound to **port 20**, and **DNS** servers are bound to **port 53**. Here is a list of the first 25 well known ports and the full list can be obtained at <http://www.iana.org/assignments/port-numbers>.

Keyword	Decimal	Description	References
-----	-----	-----	-----
	0/tcp	Reserved	
	0/udp	Reserved	
#		Jon Postel <postel@isi.edu>	
tcpmux	1/tcp	TCP Port Service Multiplexer	
tcpmux	1/udp	TCP Port Service Multiplexer	
#		Mark Lottor <MKL@nisc.sri.com>	
compressnet	2/tcp	Management Utility	
compressnet	2/udp	Management Utility	
compressnet	3/tcp	Compression Process	
compressnet	3/udp	Compression Process	
#		Bernie Volz <volz@cisco.com>	
#	4/tcp	Unassigned	
#	4/udp	Unassigned	
rje	5/tcp	Remote Job Entry	
rje	5/udp	Remote Job Entry	
#		Jon Postel <postel@isi.edu>	
#	6/tcp	Unassigned	
#	6/udp	Unassigned	
echo	7/tcp	Echo	
echo	7/udp	Echo	
#		Jon Postel <postel@isi.edu>	
#	8/tcp	Unassigned	
#	8/udp	Unassigned	
discard	9/tcp	Discard	
discard	9/udp	Discard	
#		Jon Postel <postel@isi.edu>	
discard	9/dccp	Discard SC:DISC	


```

437 # IETF dccp WG, Eddie Kohler <kohler@cs.ucla.edu>,
438 [RFC4340]
439 # 10/tcp Unassigned
440 # 10/udp Unassigned
441 sysstat 11/tcp Active Users
442 sysstat 11/udp Active Users
443 # Jon Postel <postel@isi.edu>
444 # 12/tcp Unassigned
445 # 12/udp Unassigned
446 daytime 13/tcp Daytime (RFC 867)
447 daytime 13/udp Daytime (RFC 867)
448 # Jon Postel <postel@isi.edu>
449 # 14/tcp Unassigned
450 # 14/udp Unassigned
451 # 15/tcp Unassigned [was netstat]
452 # 15/udp Unassigned
453 # 16/tcp Unassigned
454 # 16/udp Unassigned
455 qotd 17/tcp Quote of the Day
456 qotd 17/udp Quote of the Day
457 # Jon Postel <postel@isi.edu>
458 msp 18/tcp Message Send Protocol
459 msp 18/udp Message Send Protocol
460 # Rina Nethaniel <---none--->
461 chargen 19/tcp Character Generator
462 chargen 19/udp Character Generator
463 ftp-data 20/tcp File Transfer [Default Data]
464 ftp-data 20/udp File Transfer [Default Data]
465 ftp 21/tcp File Transfer [Control]
466 ftp 21/udp File Transfer [Control]
467 # Jon Postel <postel@isi.edu>
468 ssh 22/tcp SSH Remote Login Protocol
469 ssh 22/udp SSH Remote Login Protocol
470 # Tatu Ylonen <ylo@cs.hut.fi>
471 telnet 23/tcp Telnet
472 telnet 23/udp Telnet
473 # Jon Postel <postel@isi.edu>
474 24/tcp any private mail system
475 24/udp any private mail system
476 # Rick Adams <rick@UUNET.UU.NET>
477 smtp 25/tcp Simple Mail Transfer
478 smtp 25/udp Simple Mail Transfer

```

```

479 12.3 When one computer on the network wants to make use of a specific
480 service that is running on another computer on the network, the first
481 computer creates a message, places the port number of the desired service
482 into the message, and then sends it to the destination computer. If a
483 process that implements the well known service for that port is bound to the
484 port, then this process will receive the message and perform the requested
485 work.

```

486 12.4 The main restriction on **processes** that are bound to ports in the well
487 known ports range is that they **must** be running with **superuser**
488 **privileges**.

489 12.5 Registered ports (1024 - 49151)

490 12.5.1 **Registered ports** work similarly to **well known ports** except that
491 the **processes** that are bound to them **do not** need to be running with
492 **superuser privileges**. The list of **registered ports** is included in the
493 same **IANA document** that contains the list of **well known ports**.

494 12.6 Dynamic/private ports (49152 - 65535)

495 12.6.1 These ports are used as needed, and they do not have any specific
496 type of process associated with them. A typical use of the ports in this
497 range is for a web browser to make an outgoing connection with a web
498 server.

499 13 The SSH (Secure SHell) service

500 13.1 In the “Installing Gentoo Linux” book, we discussed what system services
501 were, and then we installed a **logging** service and a **cron** service. These
502 two services are accessed through software calls, but **some system**
503 **services are bound to well known ports and they make their services**
504 **available through these ports**. An example of a service that makes itself
505 available through a well known port is the **SSH** (Secure SHell) service and it
506 is usually bound to port **22**.

507 13.2 The **SSH service** allows a person to log into one computer on a network
508 from another computer on the network. The person must know the
509 **username** and **password** for an account on the remote machine before
510 logging into it, and the remote machine must have a SSH service (in the
511 form of a process) running and bound to port 22. SSH is able to provide a
512 secure connection between the machines by encrypting the data that is
513 passed between them.

514 13.3 When a system service is emerged, it usually places a small program (or
515 script) in the **/etc/init.d** directory which will allow it to be **started, stopped**
516 **and restarted**. Lets look into this directory to see what system service
517 control scripts it contains:

518 `kosan1 / # cd /etc/init.d`

```
519 kosan1 init.d # pwd
520 /etc/init.d
```

```
521 kosan1 init.d # ls
522 bootmisc      depscan.sh      hdparm          net.enp0s3      rmnologin       sshd
523 checkfs        dhcpcd          hostname        net.lo          rsyncd          syslog-ng
524 checkroot      dhcrelay        keymaps         netmount        runscript.sh    urandom
525 clock          functions.sh    local           nscd            shutdown.sh     vixie-cron
526 consolefont    gpm            localmount      numlock         slapd           xdm
527 crypto-loop    halt.sh         modules         reboot.sh       slurpd
```

528 13.4 There are a number of system service control scripts in this directory.
 529 Notice that scripts for starting, stopping and restarting the **net.np0s3**
 530 networking service, the **syslog-ng** service and the **vixie-cron** service are
 531 present here along the the script for controlling the **sshd** service. We
 532 needed to emerge the **syslog-ng** and the **vixie-cron** services ourselves but
 533 the **sshd** service is so popular that the Gentoo developers installed it by
 534 default.

535 13.5 If you want to see which of these services are currently running on your
 536 machine, you can execute the **rc-status** command:

```
537 kosan1 init.d # rc-status
538 Runlevel: default
539 net.enp0s3          [ started ]
540 syslog-ng          [ started ]
541 vixie-cron          [ started ]
542 netmount            [ started ]
543 local              [ started ]
544 Dynamic Runlevel: hotplugged
545 Dynamic Runlevel: needed
546 Dynamic Runlevel: manual
547 sshd                [ stopped ]
```

548 13.6 If you want to **start** the **sshd** service, execute the **sshd** script and pass it
 549 a **start** option. After the **sshd** service has been started, execute the
 550 **rc-status** command again to verify that it has indeed been started:

```
551 kosan1 init.d # /etc/init.d/sshd start
552 * Starting sshd ... [ ok ]

553 kosan1 init.d # rc-status
554 Runlevel: default
555 net.enp0s3          [ started ]
556 syslog-ng          [ started ]
557 vixie-cron          [ started ]
558 netmount            [ started ]
559 local              [ started ]
560 Dynamic Runlevel: hotplugged
561 Dynamic Runlevel: needed
```

```
562 Dynamic Runlevel: manual
563 sshd
```

```
[ started ]
```

564 13.7 If you want to **stop** the service, pass the **stop** option to it and if you want
565 to **restart** it, use the **restart** option. Finally, if you want the **sshd service**
566 to be automatically started when the system enters the **default runlevel**
567 (which was discussed in the “Installing Gentoo Linux” book) use the
568 **rc-update** command:

```
569 kosan1 init.d # rc-update add sshd default
570 * sshd added to runlevel default
```

571 14 Using SSH to remotely log into a machine

572 14.1 Now that you know how to start and stop services, lets use the SSH
573 service to **remotely log into a Gentoo Linux machine**. If you are working
574 with a friend, have them create an account for you on their machine and
575 then try to remotely log into it from your system using the **ssh client**
576 **program**. Also, you can ssh from your machine's normal operating system
577 into your VirtualBox Gentoo system. If you are running Windows, you can
578 us the **putty.exe** program for this (see below). **(Note: make sure your**
579 **VirtualBox network connection is set to "bridged" in the VirtualBox**
580 **settings):**

```
581 tkosan@kosan_laptop / $ ssh tkosan@206.21.94.136
582 The authenticity of host '206.21.94.136 (206.21.94.136)' can't be established.
583 RSA key fingerprint is f6:2b:63:33:99:6c:57:37:b4:b7:2d:ba:bc:da:be:77.
584 Are you sure you want to continue connecting (yes/no)? yes
585 Warning: Permanently added '206.21.94.136' (RSA) to the list of known hosts.
586 Password:
587 Last login: Sun Mar  4 15:38:33 2007 from 10.0.1.3

588 tkosan@kosan1 ~ $ pwd
589 /home/tkosan
```

590 14.2 The machine that I logged into above was at IP address **206.21.94.136**
591 and the **username** for my account on that system is **tkosan**. The
592 **tkosan@206.21.94.136** data that I passed to the **ssh** program told it that I
593 wanted to have it remotely log into the **tkosan** account on the machine on
594 the network that has IP address **206.21.94.136**.

595 14.3 When the **ssh** client program is asked to log into a remote machine for
596 the first time, it tells the user that it does not currently have encryption
597 information for this host, and it asks if it should continue. Answer by typing
598 the word **"yes"**. The program then indicates that it added information about
599 this host to a known hosts list, and it will not ask the question again in the

600 future.

601 14.4 The **known hosts** list is contained in the user's home directory in a file
602 called **.ssh/known_hosts** (notice that the **.ssh** directory is a hidden
603 directory). If you experience trouble using **ssh** to log into a computer in the
604 future, you may need to edit the **known_hosts** file or delete it so that **ssh**
605 can regenerate it from scratch.

606 14.5 In the above listing, the **hostname** of the machine I am logging in from
607 is **kosan_laptop** and the **hostname** of the machine I am logging into is
608 **kosan1**. The user account I am using on both machines is **tkosan**.

609 14.6 If you do not have a second Gentoo Linux machine available to
610 experiment with, you can download a program called **putty.exe** that you can
611 install on a Windows machine, and it will allow you to remotely log into a
612 Gentoo machine that is running the **sshd** service. The **putty.exe** program
613 can be downloaded from
614 (<http://www.chiark.greenend.org.uk/~sgtatham/putty/download.html>).

615 14.7 When you are done using **ssh** to remotely log into a machine, execute an
616 **exit** command to close the session.

617 15 Using **scp** to copy files between machines on the network

618 15.1 The **SSH** service is not only able to allow a user to log into a remote
619 machine, it can also be used to copy files between machines on the network.
620 The Linux client program for copying files is called **scp** (Secure Copy) and a
621 popular **Windows scp** client, called **pscp.exe**, can be obtained from the
622 same url that **putty.exe** was obtained from.

623 15.2 We can experiment with **scp** by copying a file from the local machine to a
624 remote machine and then copying a separate file from the remote machine
625 to the local machine. First, change into your home directory on the **local**
626 machine and **create a file called localfile.txt using a text editor** (just
627 place a line or two of text in this file). Then use the **scp** command (or the
628 **pscp** command if you are on a Windows machine) as shown below to copy
629 the file from the **local** machine to your account on the **remote** machine:

```
630 tkosan@kosan1 ~ $ scp localfile.txt tkosan@206.21.94.136:
```

```
631 Password:
```

```
632 localfile.txt          100%   16    0.0KB/s   00:00
```

633 15.3 Verify that the file was copied to the remote machine by using **ssh** (or
634 **putty**) to log into it. The **first parameter** after the **scp** command indicates

635 where the file is being **copied from** and the **second parameter** indicates
636 where it is being **copied to**. If an IP address is not present in either the
637 source or destination parameter, that means that the local machine is being
638 referenced. The colon ':' that is placed after the IP address indicates that
639 the file was copied into the user's home directory, and this is where you
640 should look for the file. You can also add a **path** after the colon if you want
641 to copy the file to a directory other than the user's home directory.

642 15.4 Now, create a file on the remote machine called **remotefile.txt**, exit
643 back to the **local** machine, and then execute the following **scp** command to
644 copy the **remotefile.txt** file from the **remote** machine to the **local** machine:

```
645 tkosan@kosan1 ~ $ scp tkosan@206.21.94.136:remotefile.txt .  
646 Password:  
647 remotefile.txt                                100%    6    0.0KB/s   00:00
```

648 15.5 In this case, the **source** is the user's home directory on the **remote**
649 machine and the **destination** is the current directory on the **local** machine.
650 If you recall from the "Installing Gentoo Linux" book, a period '.' indicates
651 the current directory.